

Computational Statistics II

Unit A.1: R programming and MCMC

Tommaso Rigon

University of Milano-Bicocca

Ph.D. in Economics and Statistics

Unit A.1

Main concepts

- Brief recap of inference via MCMC and the Metropolis–Hastings algorithm
 - Weibull model with censored data
 - Writing clean and efficient **R** code
 - Reparametrizations and transformations.
-
- Associated **R** code: https://tommasorigon.github.io/CompStat/exe/un_A1.html

Bayesian inference (recap)

- Let \mathbf{X} be the data, following some distribution $\pi(\mathbf{X} | \theta)$, i.e. the **likelihood**, with $\theta \in \Theta \subseteq \mathbb{R}^p$ being an unknown set of parameters.
- Let $\pi(\theta)$ be the **prior distribution** associated to θ .
- In Bayesian analysis, inference is based on the **posterior distribution** for θ , defined as

$$\pi(\theta | \mathbf{X}) = \frac{\pi(\theta)\pi(\mathbf{X} | \theta)}{\int_{\Theta} \pi(\theta)\pi(\mathbf{X} | \theta)d\theta}.$$

- **Key issue**: the **normalizing constant**, i.e. the above integral, is often **intractable** \implies no analytical solutions.
- Numerical approximations of $\int_{\Theta} \pi(\theta)\pi(\mathbf{X} | \theta)d\theta$ are highly unstable, especially in high dimensions \implies the **integrate R** function will not work in most cases.

Inference via MCMC (recap)

- **Key solution.** It is possible to **sample** from the posterior distribution even without knowing the normalizing constant.
- If we can get **random samples** $\theta^{(1)}, \dots, \theta^{(R)}$ from the posterior distribution, then we can approximate any functional of interest, i.e.

$$\mathbb{E}(f(\theta) \mid \mathbf{x}) \approx \frac{1}{R} \sum_{r=1}^R f(\theta^{(r)}).$$

- This approximation is often justified by the **ergodic theorem**.
- The samples $\theta^{(1)}, \dots, \theta^{(R)}$ are often dependent and may follow a Markov Chain \implies Markov Chain Monte Carlo (MCMC).

Metropolis-Hastings algorithm (recap)

- A simple strategy for posterior sampling is the **Metropolis-Hastings** (MH) algorithm.
- Set the first value of the chain θ_1 to some (reasonable) value.

At the r th value of the chain

- Let $\theta = \theta^{(r)}$ be the current status of the chain.
- Sample θ^* from a **proposal distribution** $q(\theta^* | \theta)$.
- Compute the **acceptance probability**, defined as

$$\alpha = \min \left\{ 1, \frac{\pi(\theta^* | \mathbf{X})}{\pi(\theta | \mathbf{X})} \frac{q(\theta | \theta^*)}{q(\theta^* | \theta)} \right\} = \min \left\{ 1, \frac{\pi(\theta^*)\pi(\mathbf{X} | \theta^*)}{\pi(\theta)\pi(\mathbf{X} | \theta)} \frac{q(\theta | \theta^*)}{q(\theta^* | \theta)} \right\}.$$

- With probability α , **update the status** of the chain and set $\theta \leftarrow \theta^*$.

Implementation of MCMC

- The MH is perhaps the simplest MCMC algorithm and it has several limitations. We will discuss modifications / extensions of the MH in the next units.
- In units A.1-A.2 we will focus on **practical considerations** concerning the implementation with **R** and its interface with C++ using **Rcpp**.
- This is far from a comprehensive guide about **R** programming. We will consider a specific model and we will implement the relevant code in **R**.

What about BUGS / JAGS / Stan?

- If the performance is not a concern, Stan-like software is an extremely useful tool for **practitioners** who wish to implement standard Bayesian models.
- Conversely, any non-standard or novel model, i.e. those usually developed by researchers in statistics, may be difficult or even impossible to implement.
- Besides, the “manual” implementation is very useful to **gain insights** about the model itself and it facilitates a lot the **debugging** process.

Example: Weibull model for censored data

- We consider an example from survival analysis, i.e. the data are **survival times** which may be **censored**.
- In this example, we assume that the survival times are iid random variable following a Weibull distribution $\text{Weib}(\alpha, \beta)$.
- The observed survival time t_i is either **complete** ($d_i = 1$) or **right censored** ($d_i = 0$), meaning that the survival time is higher than the observed t_i .
- The **hazard** and **survival** functions of a Weibull distribution are

$$h(t \mid \alpha, \beta) = \frac{\alpha}{\beta} \left(\frac{t}{\beta} \right)^{\alpha-1}, \quad S(t \mid \alpha, \beta) = \exp \left\{ - \left(\frac{t}{\beta} \right)^{\alpha} \right\}.$$

- Recall that the **density** function is obtained as $f(t \mid \alpha, \beta) = h(t \mid \alpha, \beta)S(t \mid \alpha, \beta)$

Likelihood function

- The **likelihood** for this parametric model, under suitable censorship assumptions, is **proportional** to the following quantity

$$\pi(\mathbf{t}, \mathbf{d} \mid \boldsymbol{\theta}) \propto \prod_{i=1}^n h(t_i \mid \alpha, \beta)^{d_i} S(t_i \mid \alpha, \beta) = \prod_{i:d_i=1} f(t_i \mid \alpha, \beta) \prod_{i:d_i=0} S(t_i \mid \alpha, \beta),$$

with (α, β) being the parameter vector.

- **Remark** When performing (Bayesian) inference, note that the likelihood is always defined up to an **irrelevant normalizing constant** not depending on the parameters $\boldsymbol{\theta}$.
- These irrelevant constants **can and should be omitted** when performing computations, especially if they are expensive to evaluate.

Bad implementation I (use the log-scale)

- In our experiments, we make use the stanford2 dataset of the survival package.
- In first place, we need to implement the log-likelihood function, say loglik.
- The following implementation of the log-likelihood is correct but **numerically unstable**.

```
loglik_inaccurate <- function(t, d, alpha, beta) {  
  hazard <- prod((alpha / beta * (t / beta)^(alpha - 1))^d)  
  survival <- prod(exp(-(t / beta)^alpha))  
  log(hazard * survival)  
}
```

```
# Evaluate the log-likelihood at the point (0.5, 1000)  
loglik_inaccurate(t, d, alpha = 0.5, beta = 1000)  
# [1] -Inf
```

- The product of several terms close to 0 leads to numerical inaccuracies \Rightarrow **use the log-scale** instead.

Bad implementation II (initialize the output)

- This second coding attempt relies on the log-scale and is indeed numerically much more stable than the previous version.
- However, this implementation is inefficient \implies **do not increase objects' dimension.**

```
loglik_inefficient2 <- function(t, d, alpha, beta) {  
  n <- length(t) # Sample size  
  log_hazards <- NULL  
  log_survivals <- NULL  
  
  for (i in 1:n) {  
    log_hazards <- c(log_hazards, d[i] * ((alpha - 1) * log(t[i] / beta) + log(alpha / beta)))  
    log_survivals <- c(log_survivals, -(t[i] / beta)^alpha)  
  }  
  sum(log_hazards) + sum(log_survivals)  
}  
  
# Evaluate the log-likelihood at the point (0.5, 1000)  
loglik_inefficient2(t, d, alpha = 0.5, beta = 1000)  
# [1] -873.3299
```

Bad implementation III (avoid for loops)

- This third attempt avoids the previous pitfalls but it is still quite inefficient \Rightarrow **use vectorized code** whenever possible.

```
loglik_inefficient1 <- function(t, d, alpha, beta) {  
  n <- length(t) # Sample size  
  log_hazards <- numeric(n)  
  log_survivals <- numeric(n)  
  
  for (i in 1:n) {  
    log_hazards[i] <- d[i] * ((alpha - 1) * log(t[i] / beta) + log(alpha / beta))  
    log_survivals[i] <- -(t[i] / beta)^alpha  
  }  
  sum(log_hazards) + sum(log_survivals)  
}  
  
# Evaluate the log-likelihood at the point (0.5, 1000)  
loglik_inefficient1(t, d, alpha = 0.5, beta = 1000)  
# [1] -873.3299
```

Good implementation

- The following version is both **numerically stable** and **efficient**.

```
loglik <- function(t, d, alpha, beta) {  
  log_hazard <- sum(d * ((alpha - 1) * log(t / beta) + log(alpha / beta)))  
  log_survival <- sum(-(t / beta)^alpha)  
  log_hazard + log_survival  
}  
  
# Evaluate the log-likelihood at the point (0.5, 1000)  
loglik(t, d, alpha = 0.5, beta = 1000)  
# [1] -873.3299
```

- All these versions of `loglik` run in fractions of seconds. However, the `loglik` function must be executed i.e. $\sim 10^5$ times within a MH algorithm.
- Moreover, in more complex models several instances of these inefficiencies add up.

Benchmarking the code

- To understand which function works better, you need to **test its performance**.
- There exist specialized packages to do so, i.e. **R** `rbenchmark` or `microbenchmark`.
- These packages execute the code several times and report the **average execution time**.
- The column “elapsed” refer to the overall time (in seconds) over 1000 replications.

```
library(rbenchmark) # Library for performing benchmarking
```

```
benchmark(  
  loglik1 = loglik(t, d, alpha = 0.5, beta = 1000),  
  loglik2 = loglik_inefficient1(t, d, alpha = 0.5, beta = 1000),  
  loglik3 = loglik_inefficient2(t, d, alpha = 0.5, beta = 1000),  
  columns = c("test", "replications", "elapsed", "relative"),  
  replications = 1000  
)
```

```
#      test replications elapsed relative  
#1 loglik1          1000   0.014    1.000  
#2 loglik2          1000   0.079    5.643  
#3 loglik3          1000   0.412   29.429
```

A matter of style

- **Formatting your code** properly is a healthy programming practice.
- You can refer to <https://style.tidyverse.org> for a comprehensive overview of good practices in **R**.
- Quoting the **tidyverse style guide**: “Good coding style is like correct punctuation: you can manage without it, but it sure makes things easier to read”.
- The **styler** **R** package automatically restyles your code for you and it is integrated within **RStudio** as an add-in.

Good

x <- 5

Bad

x = 5

Reparametrizations I

- When performing (Bayesian) inference, the choice of the **parametrization** has strong impacts on computations.
- **General advice**: perform computations on the most convenient parametrization and then transform back the obtained samples.
- As a rule of thumb, you should use parametrizations with **unbounded domains**. This facilitates the choice of proposal distributions and could also **improve the mixing**.
- In our model, the two parameters α, β are strictly positive. Hence, a common strategy is to consider their logarithm, i.e. $\theta = (\theta_1, \theta_2) = (\log \alpha, \log \beta)$.

To log or not to log?

Roberts, G. O. and Rosenthal, J. S. (2009). *Examples of adaptive MCMC*. Journal of Computational and Graphical Statistics, **18**(2), 349–367.

Reparametrizations & priors

- When reparametrizations are involved, there two possible modeling strategies.
- Choose the prior **before** the reparametrization. In our setting, we could let for example

$$\alpha \sim \text{Ga}(0.1, 0.1), \quad \beta \sim \text{Ga}(0.1, 0.1).$$

If you do so, remember to include the **jacobian** of the transformation when considering the transformed posterior!

- Choose the prior **after** the reparametrization. In our setting, we could let for example

$$\theta_1 = \log(\alpha) \sim \text{N}(0, 100), \quad \theta_2 = \log(\beta) \sim \text{N}(0, 100).$$

This strategy is simpler as it avoids the extra step of computing the jacobian.

```
logprior <- function(theta) {  
  sum(dnorm(theta, 0, sqrt(100), log = TRUE))  
}  
  
logpost <- function(t, d, theta) {  
  loglik(t, d, exp(theta[1]), exp(theta[2])) + logprior(theta)  
}
```

The MH implementation

- Since the space of θ is unbounded, it is reasonable to select a **Gaussian random walk** as proposal distribution, namely

$$(\theta^* \mid \theta) \sim N_2(\theta, 0.25^2 I_2).$$

The choice of the variance will be discussed in unit B.1.

- Gaussian random walks are **symmetric** proposals distributions, implying that

$$q(\theta \mid \theta^*) = q(\theta^* \mid \theta),$$

which means that their ratio can be simplified ($= 1$) when computing the acceptance probability α .

- As before, make sure you compute α using the log-scale to avoid numerical instabilities.
- **Remark.** Unfortunately, there is no way to avoid for loops, which are highly inefficient \implies this justifies the usage of **Rcpp** and **RcppArmadillo**.

Metropolis-Hastings code

```
RMH <- function(R, burn_in, t, d) {  
  out <- matrix(0, R, 2) # Initialize an empty matrix to store the values  
  theta <- c(0, 0) # Initial values  
  logp <- logpost(t, d, theta) # Log-posterior  
  for (r in 1:(burn_in + R)) {  
    theta_new <- rnorm(2, mean = theta, sd = 0.25) # Propose a new value  
    logp_new <- logpost(t, d, theta_new)  
    alpha <- min(1, exp(logp_new - logp))  
    if (runif(1) < alpha) {  
      theta <- theta_new; logp <- logp_new # Accept the value  
    }  
    if (r > burn_in) {  
      out[r - burn_in, ] <- theta # Store the values after the burn-in period  
    }  
  }  
  out  
}
```

```
# Executing the code  
library(tictoc) # Library for "timing" the functions  
tic()  
fit_MCMC <- RMH(R = 50000, burn_in = 5000, t, d)  
toc()  
# 0.92 sec elapsed
```

Estimated survival function

- Posterior mean of the survival function with pointwise 95% credible intervals.

